



Slides and notebooks: <https://ml4ns.github.io>

Notes by: *Nan Fletcher-Lloyd, Payam Barnaghi*

## Lecture 7. Convolutional Neural Networks (CNNs)

In lecture 7, we cover the basic principles of Convolutional Neural Networks (CNNs). Expanding on our previous lecture on neural networks and multi-layer perceptrons, here we introduce convolutional neural networks (CNNs), a specialised kind of neural network for processing data that has a known grid-like topology, such as image data (a 2D grid of pixels) and time-series data (a 1D grid when samples are taken at regular time intervals).

We further expand on the structure of a CNN, including a description of the convolution and pooling processes. In a convolutional layer, a sliding kernel is applied to an input to produce an output. In this lecture, we consider how to design our convolutional layers, including choice of padding, choice of stride, and the number of input and output channels.

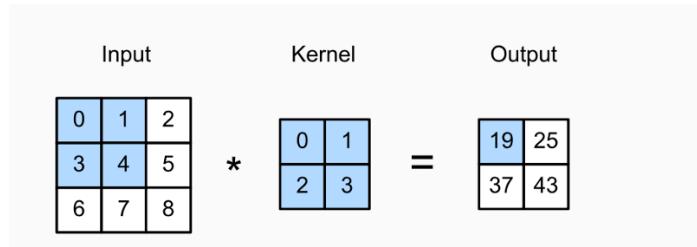


Figure 7.1. Example of a convolution  
Source: Dive into Deep Learning, Aston Zhang *et al.*, <https://d2l.ai/>

Since we typically use small kernels in a given convolution, we risk losing pixels along the image perimeter. One solution to this problem is to augment/pad our input along the boundary, thereby retaining the original input's dimensions. This is known as the same padding. If we were to apply a kernel without padding, we would produce an output of the exact dimensions as the kernel (as in Figure 7.1), and this is known as valid padding. Figure 7.1 illustrates a convolution that might occur in a CNN by taking an input and applying a sliding kernel (convolution window).

The pooling layer of a CNN is responsible for reducing the spatial size of the previous convolved layer. This is to decrease the computational power required to process the data through dimensionality reduction. Here, we introduce two types of pooling that can occur in a pooling layer: maximum pooling returns the maximum value from the portion of the input covered by the kernel, and average pooling returns the average of all values in that portion. Figure 7.2 illustrates a  $2 \times 2$  Maximum pooling

Input	Kernel	Output													
<table border="1"> <tr><td>0</td><td>1</td><td>2</td></tr> <tr><td>3</td><td>4</td><td>5</td></tr> <tr><td>6</td><td>7</td><td>8</td></tr> </table>	0	1	2	3	4	5	6	7	8	*	<table border="1"> <tr><td>0</td><td>1</td></tr> <tr><td>2</td><td>3</td></tr> </table>	0	1	2	3
0	1	2													
3	4	5													
6	7	8													
0	1														
2	3														

Figure 7.2. Example of maximum pooling  
Source: Dive into Deep Learning, Aston Zhang et al., <https://d2l.ai/>

Having successfully enabled the model to understand the features, we flatten the final output and feed it to a regular neural network for classification (e.g., *Softmax*). A fully connected layer is often added to learn non-linear combinations of the high-level features as represented by the output of the convolution layer.

Finally, we discuss the various architectures of CNNs that have been key to solving current machine learning and decision-making problems. The lab notebook/practice aims to show how to work with *PyTorch* libraries and to demonstrate the key techniques needed to build a CNN for the classification of image data.