

IMPERIAL

Ensemble Models and Kernel-based Methods

Payam Barnaghi
Department of Brain Sciences &
School of Convergence Science (Human and Artificial Intelligence)
Imperial College London
January 2026



1

IMPERIAL

Feature Vectors

- So far, we have assumed that each item we wish to classify, cluster, or process can be represented as a fixed-size feature vector.
- However, for certain types of items/concepts, it is not clear how to best represent them as fixed-sized feature vectors.
- For example, how do we represent a text document or protein sequence, which can be of variable length? Or a molecular structure, which has a complex 3d geometry? or an evolutionary tree, which has variable size and shape?

2

2

Kernel functions

IMPERIAL

- One approach to the latter problem is to assume that we have a way to measure the similarity between items that doesn't require preprocessing them into a feature vector.
- For example, we can apply a function to transform the original data into a higher-dimensional space in which the data items can then be compared/processed.
- This process can be done via kernel functions.

3

3

Linear separability

IMPERIAL


- We have also often assumed that the data items are linearly separable, but what if the data were not linearly separable?
- In this case, we can again apply a kernel that maps each data instance (from the original non-linear observation space) into a higher-dimensional space where the data instances become separable.

4

4

Example

IMPERIAL





Source: https://xavierbourretscotte.github.io/Kernel_feature_map.html

5

5

Example – applying polynomial kernel

IMPERIAL



Source: https://xavierbourretscotte.github.io/Kernel_feature_map.html

6

6

What is a polynomial kernel?

IMPERIAL

- If our data is a set of (x_1, x_2)
- Transform it to (z_1, z_2, z_3) in which:

$$z_1 = x_1$$

$$z_2 = x_2$$

$$z_3 = x_1^2 + x_2^2$$

$$(z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

- This is, in fact, the sum of polynomials

7

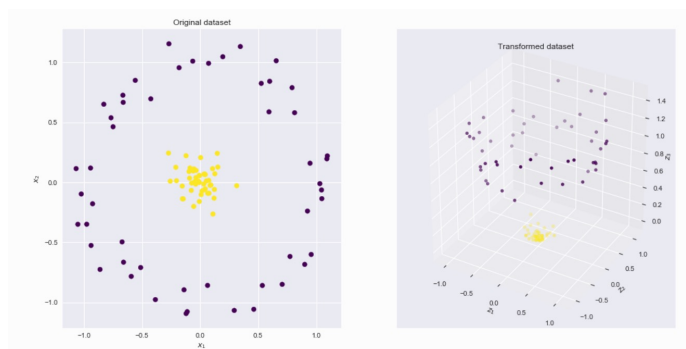
7

A polynomial kernel

IMPERIAL

$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

```
def feature_map_1(X):
    return np.asarray((X[:,0], X[:,1], X[:,0]**2 + X[:,1]**2)).T
```



Source: https://xavierbourretscotte.github.io/Kernel_feature_map.html

8

8

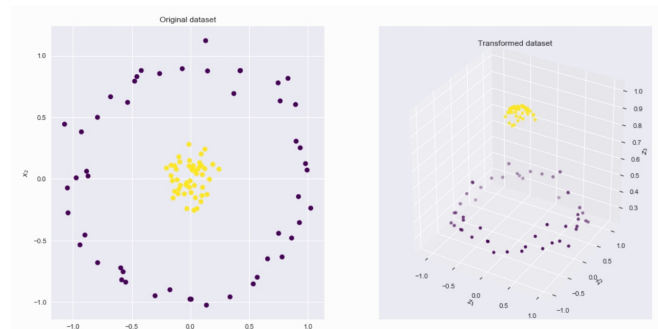
Gaussian Radial Basis Function (RBF)

IMPERIAL

Gaussian Radial Basis Function (RBF) centered at 0, 0

$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, e^{-(x_1^2 + x_2^2)})$$

```
def feature_map_2(X):
    return np.asarray((X[:,0], X[:,1], np.exp(-(X[:,0]**2 + X[:,1]**2)))).T
```



Source: https://xavierbourretsicotte.github.io/Kernel_feature_map.html

9

9

RBF Kernel in more generalised form*

IMPERIAL

The **squared exponential kernel** (SE kernel) or **Gaussian kernel** is defined by

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^T \Sigma^{-1}(\mathbf{x} - \mathbf{x}')\right)$$

If Σ is diagonal, this can be written as

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2} \sum_{j=1}^D \frac{1}{\sigma_j^2} (x_j - x'_j)^2\right)$$

Source: Kevin Murphy, Machine Learning: A probabilistic Approach, MIT Press.

10

10

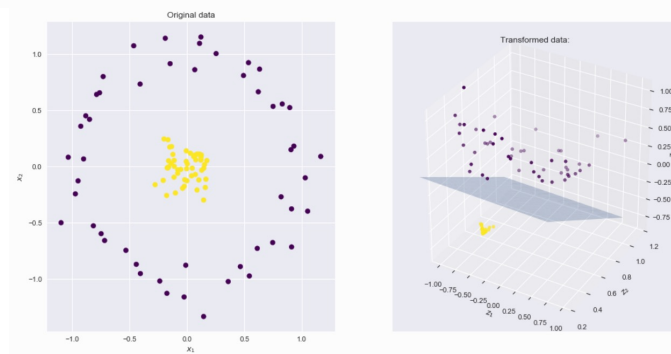
Another polynomial function

IMPERIAL

$$x_1, x_2 \mapsto z_1, z_2, z_3$$

$$z_1 = \sqrt{2}x_1x_2 \quad z_2 = x_1^2 \quad z_3 = x_2^2$$

```
def feature_map_3(X):
    return np.asarray(( np.sqrt(2) *X[:,0] * X[:,1], X[:,0]**2, X[:,1]**2)).T
```



Source: https://xavierbourretsicotte.github.io/Kernel_feature_map.html

11

11

Support Vector Machines

IMPERIAL

- Support vector machines (SVMs) are a set of supervised learning methods used for classification, regression and outlier detection.
- The advantages of support vector machines are:
 - Effective in high-dimensional spaces.
 - It could still be effective in cases where the number of dimensions is greater than the number of samples.
 - It uses a subset of training points in the decision function (called support vectors), making it also memory-efficient.
 - Versatile: Different kernel functions can be specified for the decision function. Common kernels are provided, but it is also possible to define custom kernels.

Source: Scikit learn documentation

12

12

SVM - disadvantages

IMPERIAL

- The disadvantages of support vector machines include:
 - If the number of features exceeds the number of samples, the method will likely give poor performance.
 - SVMs **do not directly provide probability estimates**.

13

13

SVM Classifier

IMPERIAL


- A Support Vector Machine (SVM) is a discriminative classifier formally defined by a separating hyperplane.
- In other words, given labelled training data (supervised learning), the algorithm outputs an optimal hyperplane which categorises new examples.
- e.g., For a linearly separable set of 2D-points which belong to one of two classes, find a separating straight line.

14

14

SVM Example

IMPERIAL



Source: https://vovkos.github.io/doxyrest-showcase/opencv/sphinx_rtd_theme/page_tutorial_introduction_to_svm.html

15

15

Linear separation

IMPERIAL

- In this example, we deal with lines and points in the Cartesian plane instead of hyperplanes and vectors in a high-dimensional space.
- This is a simplification of the problem.
- It is important to understand that this is done only because our intuition is better built from easily imagined examples.
- However, the same concepts apply to tasks where the examples to classify lie in a space whose dimension is higher than two.

Source: https://vovkos.github.io/doxyrest-showcase/opencv/sphinx_rtd_theme/page_tutorial_introduction_to_svm.html

16

16

SVM – linear separation

IMPERIAL

- We can intuitively define a criterion to estimate the worth of the lines:
 - A line is bad if it passes too close to the points because it will be noise-sensitive and will not generalise correctly. Therefore, our goal should be to find the line passing as far as possible from all points.
 - Then, the SVM algorithm operates by finding the hyperplane that maximises the minimum distance to the training examples.
 - This distance is given the important name of the margin in SVM. Therefore, the optimal separating hyperplane maximises the margin of the training data.


Source: https://vovkos.github.io/doxyrest-showcase/opencv/sphinx_rtd_theme/page_tutorial_introduction_to_svm.html

17

17

Example: SVM for a linearly separable set of 2d-points

IMPERIAL




Source: https://vovkos.github.io/doxyrest-showcase/opencv/sphinx_rtd_theme/page_tutorial_introduction_to_svm.html

18

18

Mathematical formulation*

IMPERIAL



$$\sum_{i \in SV} y_i \alpha_i K(x_i, x) + b,$$


Source: <https://scikit-learn.org/stable/modules/svm.html>

19

19

SVM with different kernels

IMPERIAL

Source: https://vovkos.github.io/doxyrest-showcase/opencv/sphinx_rtd_theme/page_tutorial_introduction_to_svm.html

20

20

Decision Trees

IMPERIAL

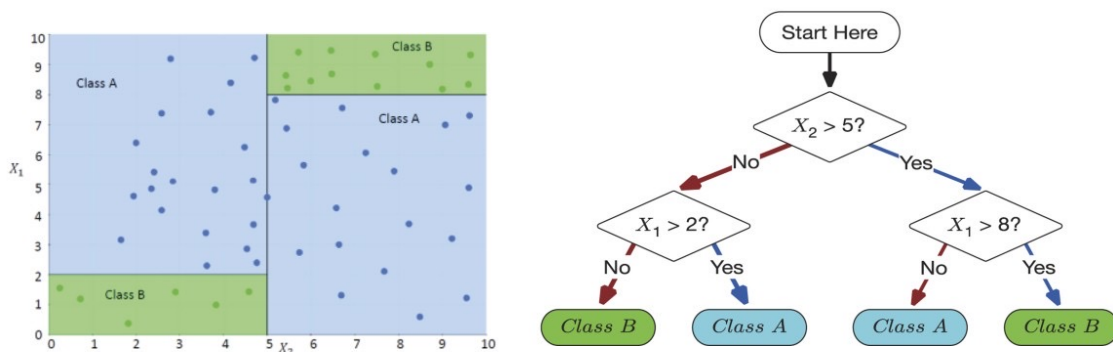
- A Decision Tree (DT) model is an interpretable model in which a series of comparisons of predictor values against threshold values determines the final output.
- DTs are made up of nodes, branches, and leaves.
- Each node represents a feature, each branch a choice, and each leaf an outcome.
- DTs take a top-down approach to data, attempting to group and classify similar observations and to search for optimal criteria to partition dissimilar observations until they reach a particular level of similarity.

21

21

Decision Making Using DTs

IMPERIAL



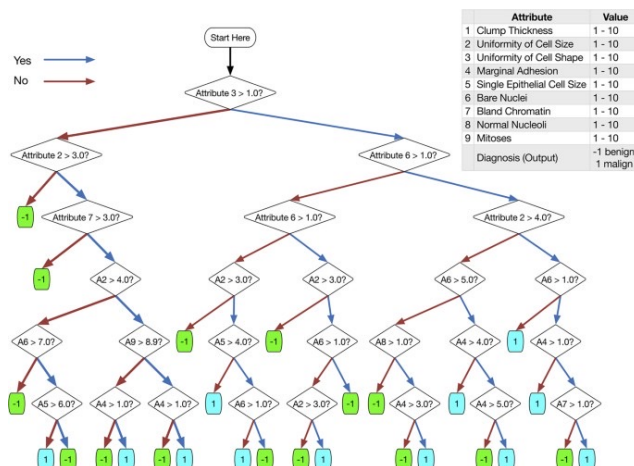
Source: Valdes, G., Luna, J., Eaton, E. et al. MediBoost: a Patient Stratification Tool for Interpretable Decision Making in the Era of Precision Medicine. *Sci Rep* 6, 37854 (2016).
<https://doi.org/10.1038/srep37854>

22

22

DTs in Clinical Applications

IMPERIAL



Example Interpretable Rules Induced by MediBoost:

A3 Uniformity of Cell Shape ≤ 1.0 \wedge A2 Uniformity of Cell Size > 3.0 \wedge A7 Bland Chromatin ≤ 3.0 \Rightarrow predict benign
 A3 Uniformity of Cell Shape > 1.0 \wedge A6 Bare Nuclei ≤ 1.0 \wedge A2 Uniformity of Cell Size ≤ 3.0 \Rightarrow predict benign

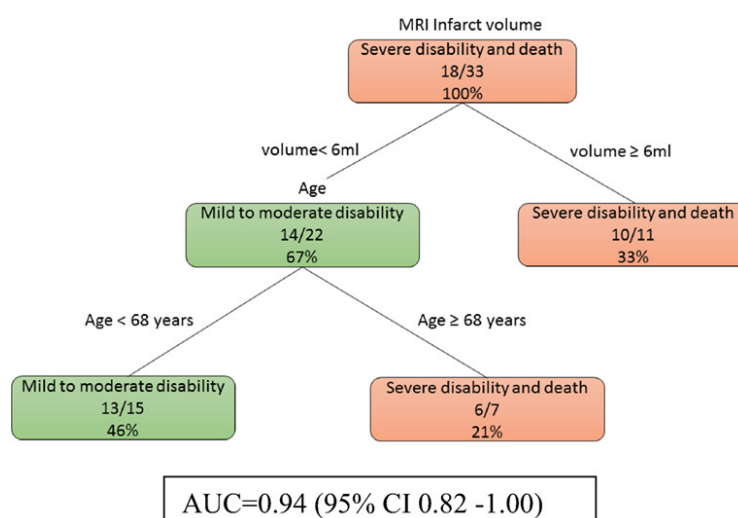
Source: Valdes, G., Luna, J., Eaton, E. et al. MediBoost: a Patient Stratification Tool for Interpretable Decision Making in the Era of Precision Medicine. *Sci Rep* 6, 37854 (2016).
<https://doi.org/10.1038/srep37854>

23

23

A decision tree model in TBI

IMPERIAL



Source: Phan, T. G., Chen, J., Singhal, S., Ma, H., Clissold, B. B., Ly, J., & Beare, R. (2018). Exploratory Use of Decision Tree Analysis in Classification of Outcome in Hypoxic-Ischemic Brain Injury. *Frontiers in Neurology*, 9. <https://doi.org/10.3389/fneur.2018.00126>

24

24

Learning decision trees

IMPERIAL

- Decision Trees (DTs) are a non-parametric supervised learning method used for classification and regression.
- The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features.
- A tree can be seen as a piecewise constant approximation.

Source: <https://scikit-learn.org/stable/modules/tree.html>

25

25

Learning decision trees

IMPERIAL

- There are different algorithms that can be used to learn decision tree structure from the data.
- In principle, these algorithms start with a root node in the tree structure, explore the features, and create rules that separate the data one step at a time.
- One of the basic algorithms is ID3.

26

26

The ID3 algorithms

IMPERIAL

- Start with the dataset and set the root node.
- Start iterating through the dataset, and at each step, go through the features of the data and calculate the Entropy(\mathcal{H}) and Information gain (IG) of the features.
- Select the feature with the highest information gain, which is equivalent to the feature that reduces entropy the most.
- Split the tree by the selected feature to create a new branch in the tree structure.
- Continue the algorithm by considering only the features that have never been selected before.

27

27

Entropy and Information Gain*

IMPERIAL

Consider a dataset with $S = \{9 \text{ positive}, 5 \text{ negative}\}$.

Step 1: Compute Entropy of S :

$$H(S) = - \left(\frac{9}{14} \log_2 \frac{9}{14} + \frac{5}{14} \log_2 \frac{5}{14} \right)$$

$$H(S) \approx 0.940$$

Step 2: Split on a Feature A :

Suppose splitting S based on A results in two subsets:

- $S_1 = \{6 \text{ positive}, 2 \text{ negative}\}$,
- $S_2 = \{3 \text{ positive}, 3 \text{ negative}\}$.

Calculate the weighted entropy after the split:

$$H(S|A) = \frac{8}{14} H(S_1) + \frac{6}{14} H(S_2)$$

$$H(S|A) = \frac{8}{14} \left(-\frac{6}{8} \log_2 \frac{6}{8} - \frac{2}{8} \log_2 \frac{2}{8} \right) + \frac{6}{14} \left(-\frac{3}{6} \log_2 \frac{3}{6} - \frac{3}{6} \log_2 \frac{3}{6} \right)$$

$$H(S|A) \approx 0.788$$

Step 3: Compute Information Gain:

$$IG(S, A) = H(S) - H(S|A)$$

$$IG(S, A) = 0.940 - 0.788 = 0.152$$

The feature A has an Information Gain of 0.152.

28

28

Decision Trees - Limitations

IMPERIAL

- Prone to overfit
- Poor generalisation performances
- High variance
 - A slight change in the data can result in a completely different set of splits, which can make interpretation difficult
- They can be inherently unstable
 - The effect of an error in the top splits propagates down to all the splits below due to their hierarchical nature
- Produce biased trees for an unbalanced dataset

29

29

How to overcome these limitations?

IMPERIAL

- Ensemble:
 - Single DT is fast but does not perform well
 - Learn from multiple trees
- We need to be careful not to learn the same tree over and over again.
- Bagging:
 - Bootstrap aggregating, a method that results in low variance
 - Construct N trees and learn a classifier for each bootstrap sample, and average them
 - Can improve the accuracy of unstable models that tend to overfit
- Feature Randomness:
 - Feature bagging generates a random subset of features, which ensures low correlation among models


30

30

Random Forest

IMPERIAL

- Random Forest was introduced in 2001 to overcome the limitations of decision trees [1].
- Ensemble model consisting of multiple decision trees.



[1] Breiman, Leo. "Random forests." Machine learning 45.1 (2001): 5-32.

31

31

Ensemble models

IMPERIAL

- Ensemble methods combine the predictions of several base estimators trained with a given learning algorithm to improve generalisability and robustness relative to a single estimator.
- Two families of ensemble methods are usually distinguished:
 - Averaging methods
 - Boosting methods

Source: <https://scikit-learn.org/stable/modules/ensemble.html>

32

32

Averaging methods

IMPERIAL

- In averaging methods, the driving principle is to build several estimators independently and then average their predictions.
- The combined estimator is usually better than any single base estimator because its variance is reduced.
- Examples: Bagging methods

Source: <https://scikit-learn.org/stable/modules/ensemble.html>

33

33

Bagging

IMPERIAL

- There are different Bagging methods, but they mostly differ from each other in the way they draw random subsets of the training set:
 - When random subsets of the dataset are drawn as random subsets of the samples, then this algorithm is known as **Pasting**.
 - When samples are drawn with replacement, then the method is known as **Bagging**.
 - When random subsets of the dataset are drawn as random subsets of the features, then the method is known as **Random Subspaces**.
 - Finally, when base estimators are built on subsets of both samples and features, then the method is known as **Random Patches**.
- In scikit-learn, bagging methods are offered as a unified [BaggingClassifier](#) meta-estimator (resp. [BaggingRegressor](#)), taking as input a user-specified estimator along with parameters specifying the strategy to draw random subsets.

Source: <https://scikit-learn.org/stable/modules/ensemble.html>

34

34

Boosting methods

IMPERIAL

- By contrast, in **boosting methods**, base estimators are built sequentially, and one tries to reduce the bias of the combined estimator.
- The motivation is to combine several weak models to produce a powerful ensemble.
- Examples: XGBoost, Gradient Tree Boosting

Source: <https://scikit-learn.org/stable/modules/ensemble.html>

35

35

Gradient Boosting

IMPERIAL

- The term “gradient boosting” comes from the idea of “boosting” or improving a single weak model by combining it with several other weak models to generate a collectively strong model.
- Gradient boosting sets targeted outcomes for the next model to minimise errors.
- Targeted outcomes for each case are based on the gradient of the error (hence the name gradient boosting) with respect to the prediction.

Source: NVIDIA, <https://www.nvidia.com/en-us/glossary/data-science/xgboost/>

36

36

XGBoost

IMPERIAL

- GBDTs iteratively train an ensemble of shallow decision trees, with each iteration using the residuals from the previous model to fit the next.
- The final prediction is a weighted sum of all of the tree predictions. Random forest “bagging” minimises the variance and overfitting, while GBDT “boosting” minimises the bias and underfitting.
- With XGBoost, trees are built in parallel instead of sequentially like GBDT.
- XGBoost follows a level-wise strategy, scanning across gradient values and using these partial sums to evaluate the quality of splits at every possible split in the training set.

Source: NVIDIA, <https://www.nvidia.com/en-us/glossary/data-science/xgboost/>

37

37

Random Forests

IMPERIAL

- In random forests, each tree in the ensemble is built from a sample drawn with replacement (i.e., a bootstrap sample) from the training set.
- When splitting each node during the construction of a tree, the best split is found either from all input features or a random subset of size `max_features`.
- The purpose of these two sources of randomness is to decrease the variance of the forest estimator.


Source: <https://scikit-learn.org/stable/modules/ensemble.html>

38

38

Structure of Random Forests

IMPERIAL



Source: Malebary, S.J., Khan, Y.D. Evaluating machine learning methodologies for identification of cancer driver genes. *Sci Rep* 11, 12281 (2021). <https://doi.org/10.1038/s41598-021-91656-8>


39

39

Feature important in RF: Example

IMPERIAL

- Selecting the most important self-assessed features for predicting conversion to mild cognitive impairment with random forest and permutation-based methods:
 - Gómez-Ramírez, J., et al.(2020), <https://doi.org/10.1038/s41598-020-77296-4>



(b) Stacked bar plot with the mean absolute value of the SHAP values in the horizontal axis and the features in the vertical axis. Blue represents Class 0 or Healthy and in red is depicted Class 1 or MCI converters.

40

40

Random Forests

IMPERIAL

- Individual decision trees typically exhibit high variance and tend to overfit.
- The injected randomness in forests yields decision trees with somewhat decoupled prediction errors.
- By averaging those predictions, some errors can cancel each other out. Random forests achieve a reduced variance by combining diverse trees, sometimes at the cost of a slight increase in bias.

Source: <https://scikit-learn.org/stable/modules/ensemble.html>

41

41

Practical note on RFs

IMPERIAL

- In practice, variance reduction is often significant, yielding a better overall model.
- In contrast to the original publication, the scikit-learn implementation combines classifiers by averaging their probabilistic predictions rather than letting each classifier vote for a single class.

Source: <https://scikit-learn.org/stable/modules/ensemble.html>

42

42

SHAP

IMPERIAL


- SHAP (SHapley Additive exPlanations) is a game theoretic approach to explaining the output of any machine learning model.
- It connects optimal credit allocation with local explanations using the classic Shapley values from game theory and their related extensions.
- See: <https://github.com/slundberg/shap>

43

43

SHAP example I

IMPERIAL




44

44

SHAP example 2

IMPERIAL




45

45

SHAP example 3

IMPERIAL



Capstick, A., Palermo, F., Zakka, K. et al. Digital remote monitoring for screening and early detection of urinary tract infections. *npj Digit. Med.* 7, 11 (2024). <https://doi.org/10.1038/s41746-023-00995-5>

46

46

Open discussion

IMPERIAL

- A research team is building a predictive model to estimate the progression of Parkinson's disease using patient data (motor scores, speech patterns, imaging features).
- **Ensemble Models:**
 - How could Random Forests or Gradient Boosting improve prediction accuracy compared to a single decision tree?
 - What trade-offs exist between interpretability and performance in clinical settings?
- **Kernel-Based Methods:**
 - Could an SVM with an RBF kernel capture non-linear relationships in complex neurological data?
 - How do we choose between kernel-based methods and ensemble models for this problem?
- **Ethical & Practical Considerations:**
 - Which approach is more robust to small sample sizes often seen in neurology research?
 - How do we ensure fairness and avoid bias when using these methods on diverse patient populations?

47

47

IMPERIAL

Review questions

48

Q1 – RF vs DT

IMPERIAL

- If we have a dataset with multiple features, the samples of which vary, and the different features contribute to the overall class prediction, which of DT or RF would be a good choice?
- menti code will be provided.

49

49

Q2

IMPERIAL

- In a random forest model, when samples are drawn with replacement, what is the method known?

50

50

Q3

IMPERIAL

- In SVM models, there is a hyperparameter known as the Soft Margin of SVM (shown as C in the model implementation). The C hyperparameter introduces a penalty for each misclassified sample.
- If you set C to a large value, it will imply a small margin in support vectors. What issue do you think the small margins in support vectors could then cause?

51

51

If you have any questions

IMPERIAL

- Please feel free to arrange a meeting or email (p.barnaghi@imperial.ac.uk).
- To arrange a meeting, please email my colleague, Ms Rhiannon Kirby.
- My office: 928, Sir Michael Uren Research Hub, White City Campus.

52

52